



AlgoMentor: A Socratic Algorithm Tutor Chatbot

- Khalid Al Jundi alfatto@uwindsor.ca 110157691
- Harmit Patel patel62b@uwindsor.ca 110164115

COMP-4400 Principles of Programming Languages
Advisor: Jianguo Lu jlu@uwindsor.ca



Idea

- Algorithms tutor chatbot
- helps you to solve programming questions, specifically data structures and algorithm questions like leetcode problems.



problem motivation and use case

- Students struggle with algorithm questions
- current LLM chatbots give direct answers
- we need guided, interactive learning, helping the user step by step to solve the question
- can be used as an ai chatbot on leetcode for example to help users.

demo

- help solving the leetcode TWO SUM problem

[Description](#) | [Accepted](#) | [Editorial](#) | [Solutions](#) | [Submissions](#)

1. Two Sum

[Easy](#) | [Topics](#) | [Companies](#) | [Hint](#)

Given an array of integers `nums` and an integer `target`, return *indices of the two numbers such that they add up to* `target`.

You may assume that each input would have **exactly one solution**, and you may not use the *same* element twice.

You can return the answer in any order.

Example 1:

Input: `nums = [2,7,11,15], target = 9`

Output: `[0,1]`

Explanation: Because `nums[0] + nums[1] == 9`, we return `[0, 1]`.

Example 2:

Input: `nums = [3,2,4], target = 6`

Output: `[1,2]`

Example 3:

Input: `nums = [3,3], target = 6`

Output: `[0,1]`

Constraints:

- `2 <= nums.length <= 104`
- `-109 <= nums[i] <= 109`
- `-109 <= target <= 109`
- Only one valid answer exists.**

68.4K | 1.9K | | |

1290 Online

Socratic Algorithm Tutor

Paste a problem, then ask questions. The tutor will never give you the answer — only the next hint.

Problem

Paste a LeetCode / Codeforces problem here...

Start Session

Status

Conversation

Your message

Ask a question or describe your approach...

Send

Socratic Algorithm Tutor

Paste a problem, then ask questions. The tutor will never give you the answer — only the next hint.

Problem

You can return the answer in any order.

Example 1:

Input: `nums = [2,7,11,15]`, `target = 9`

Output: `[0,1]`

Explanation: Because `nums[0] + nums[1] == 9`, we return `[0, 1]`.

Example 2:

Input: `nums = [3,2,4]`, `target = 6`

Output: `[1,2]`

Example 3:

Input: `nums = [3,3]`, `target = 6`

Output: `[0,1]`

Constraints:

`2 <= nums.length <= 104`

`-109 <= nums[i] <= 109`

`-109 <= target <= 109`

Only one valid answer exists.

Start Session

Status

Session started. Guided hint mode | stuck level: 0 | confirmed ideas: 0 | active guesses: 0

Conversation

Your message

Ask a question or describe your approach...

Send

Status

Guided hint mode | stuck level: 0 | confirmed ideas: 2 | active guesses: 1

Conversation

can i use sorting for this problem?

No, sorting changes the order, which makes the original indices harder to recover. What kind of lookup would let you avoid reordering the array?

hashmap search

Yes, that direction makes sense because it gives fast lookup. What should the key and value represent?

array value as dictionary key and index as dictionary value

Your message

`nums[i]: index`

Send

Conversation

Yes, that direction makes sense because it gives fast lookup. What should the key and value represent?

array value as dictionary key and index as dictionary value

Yes, that direction makes sense because it gives fast lookup. What should the key and value represent?

nums[i]: index

Yes, that mapping works because it lets you find a seen number and recover its position. When should you check for the partner value: before or after storing the current number?

Your message

nums[i]: index

Send



It doesn't give you the answer, it will give you a hint you and makes you think

give me the answer

No, because we were just getting started on building the solution. What's an efficient way to iterate over the array only once?



Approach

- Small local LLM (llama3.2:3b)
- Controlled using LLM programming techniques
- Acts as a tutor, not a solver


Architecture

User → Prompt + History

→ LLM (Ollama)

→ RAG (ChromaDB + embeddings)

→ Control Logic → Response






Prompt programming

System prompt defines tutor behavior

Forces step-by-step guidance

Prevents giving full solutions



```
1 SYSTEM_PROMPT = """You are an algorithms tutor for interview-style problems.
2
3 Your job is to help the student discover the solution themselves without giving away the final answer.
4
5 Hard rules:
6 - Do not provide code, pseudocode, full steps, or the final algorithm.
7 - Do not directly name the correct data structure, pattern, or algorithm unless the student already named it first.
8 - Every reply must first respond to the student's latest message, then guide them forward.
9 - Prefer one short reply made of:
10   1. a brief acknowledgement or correction
11   2. one short reason tied to what the student just said
12   3. exactly one next question
13 - If the student is correct, say so plainly, like "Yes, that makes sense because ..."
14 - If the student is partly correct, say so plainly, like "Partly, but ..."
15 - If the student is wrong, say so plainly, like "No, because ..."
16 - Keep the reply under 45 words.
17 - Never use markdown, bullets, numbered lists, or code fences.
18 - Never repeat a question from QUESTIONS ALREADY ASKED.
19 - Build on WHAT THE STUDENT HAS ALREADY CONFIRMED.
20 - If the student asks for the full answer or code, refuse briefly and ask the smallest next guiding question instead.
21
22 Teaching style:
23 - Start broad when stuck_count is low.
24 - Become more concrete as stuck_count rises.
25 - Ask about constraints, state, invariants, examples, or what information needs to be remembered.
26 - Sound like a real back-and-forth tutor, not a hint list.
27 - Use the reference only to understand the topic and common mistakes. Do not quote it and do not leak explicit answers from it.
28
29 Good behavior:
30 - "What would make checking for a needed value faster?"
31 - "If you process numbers left to right, what information might be worth remembering?"
32 - "What should the stored information help you recover at the end?"
33
34 Bad behavior:
```



More prompt programming

```
33
34  Bad behavior:
35  - Naming the exact algorithm immediately
36  - Writing code
37  - Giving the final sequence of steps
38  - Repeating a previous hint
39  """"|
```



Multi-turn conversation

Maintains chat history

Tracks student knowledge state

Tracks how many times student was stuck and how many times

Adapts responses over time



RAG (Retrieval-Augmented Generation)

Retrieves relevant algorithm content from python file

Includes this relevant content in prompt

This improves accuracy and helps the chatbot answer in a more targeted way

Example RAG usage for TwoSum Problem



```
{
```

```
  "id": "two_sum_hashmap",
```

```
  "topic": "array",
```

```
  "text": "PROBLEM: Two Sum. TUTOR GOAL: Guide the student toward a one-pass lookup-based  
solution without naming the data structure immediately. COMMON PITFALL: Sorting changes positions  
and makes original indices harder to recover. HINT_1: What if you could check whether a needed number  
was seen before in near-constant time? HINT_2: For each number x, what partner value would complete  
the target? HINT_3: If you remember previously seen numbers, what extra information would help you  
return positions instead of values?"
```

```
}
```


Why Prompt programming + RAG instead of finetuning?

- Faster to implement
- Works with small models, fine-tuning usually needs larger, cleaner training data, while RAG can work well even with a small curated knowledge base.
- Highly controllable (changing content provided to prompt to retrieve largely impacts the chatbot's response), this is very effective with our usecase since every type of data structure/algorithm problem has a unique way



What We Did NOT Use

No fine-tuning of model weights (LoRA, QLoRA)

No agents/tool-use



Limitations

Small dataset, RAG quality depends on the quality of the content provided to the LLM

No learned adaptation since we did not finetune

May still hallucinate or drift off topic



future improvements

Expand dataset

Add agent capabilities

Run user code to check correctness and if not, give next steps



Thank You !!

Any Questions